

# COMP9414 Artificial Intelligence

## GridWorldEnv Environment User Guide

Term 2, 2025

**Version:** 1.0

**Last Updated:** July 2025

**Purpose:** This guide provides comprehensive documentation for the GridWorldEnv class used in Assignment 2

## 1 Overview

The GridWorldEnv class implements a grid world environment for reinforcement learning experiments. This guide covers installation, usage, and troubleshooting for the environment provided in `env.py`.

## 2 Environment Specifications

The environment implements an  $11 \times 11$  grid world with the following characteristics:

- **Grid Size:**  $11 \times 11$  cells
- **State Space:** Discrete 2D coordinates  $(x, y)$  where  $x, y \in \{0, 1, \dots, 10\}$
- **Action Space:** 4 discrete actions
- **Goal Position:**  $(10, 10)$  - bottom-right corner
- **Starting Position:** Random (avoiding obstacles and goal)
- **Episode Termination:** Only when goal is reached

### 2.1 Coordinate System

The environment uses a standard 2D coordinate system:

- Origin  $(0, 0)$  is at the top-left corner
- X-axis increases downward (rows)
- Y-axis increases rightward (columns)

### 2.2 Obstacles

The environment contains 10 fixed obstacles arranged in two patterns:

**L-shaped pattern** (top-left area):

- Positions:  $(2, 2)$ ,  $(2, 3)$ ,  $(2, 4)$ ,  $(3, 2)$ ,  $(4, 2)$

**Cross pattern** (centre):

- Positions: (5, 4), (5, 5), (5, 6), (4, 5), (6, 5)

## 2.3 Action Space

Action	Value	Effect
Up	0	Decrease x by 1
Down	1	Increase x by 1
Left	2	Decrease y by 1
Right	3	Increase y by 1

## 2.4 Reward Structure

Event	Reward
Reach Goal	+25
Hit Obstacle	-10
Normal Step	-1

# 3 Installation and Setup

## 3.1 Prerequisites

Ensure you have Python 3.7+ and the following packages installed:

```
1 pip install numpy pygame
```

## 3.2 File Structure

Your project directory should contain:

```
your_project/
|-- env.py                # The environment file (provided)
|-- your_solution.ipynb   # Your Jupyter notebook solution
'-- images/               # Required image directory
    |-- grid_agent.png    # Agent image
    |-- grid_goal_position.png # Goal position image
    '-- grid_obstacle.png # Obstacle image
```

# 4 API Reference

## 4.1 Class: GridWorldEnv

### 4.1.1 Constructor

```
1 GridWorldEnv(seed=None)
```

#### Parameters:

- seed (int, optional): Random seed for reproducibility

#### Example:

```
1 env = GridWorldEnv(seed=42) # Reproducible environment
```

#### 4.1.2 reset()

```
1 reset() -> Tuple[int, int]
```

Resets the environment to a new episode.

##### Returns:

- state: Initial position as tuple  $(x, y)$

##### Example:

```
1 state = env.reset() # Returns e.g., (3, 7)
```

#### 4.1.3 step(action)

```
1 step(action: int) -> Tuple[Tuple[int, int], float, bool, dict]
```

Executes one environment step.

##### Parameters:

- action: Integer in  $\{0, 1, 2, 3\}$

##### Returns:

- next\_state: New position  $(x, y)$
- reward: Reward for this step
- done: Whether episode is finished
- info: Additional information (empty dict)

##### Example:

```
1 next_state, reward, done, info = env.step(0) # Move up
```

#### 4.1.4 render()

```
1 render(delay=0.1, mode="human", episode=1,  
2         learning_type="Q-learning",  
3         availability=None, accuracy=None)
```

Displays the environment using Pygame.

##### Parameters:

- delay: Time delay between frames (seconds)
- mode: Rendering mode (unused, for compatibility)
- episode: Current episode number for display

- learning\_type: Algorithm name for display
- availability: Teacher availability (0.0-1.0)
- accuracy: Teacher accuracy (0.0-1.0)

**Example:**

```
1 env.render(delay=0.1, episode=50, learning_type="SARSA")
```

**4.1.5 close()**

```
1 close() -> None
```

Closes the Pygame window and cleans up resources.

**Example:**

```
1 env.close() # Always call when done
```

## 5 Usage Examples

### 5.1 Jupyter Notebook Setup

Since your solution must be submitted as a Jupyter notebook (.ipynb), start your notebook with:

```
1 # Import required libraries
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from tqdm import tqdm
7
8 # Import the environment
9 from env import GridWorldEnv
10
11 # Set random seed for reproducibility
12 np.random.seed(42)
13
14 # Configure plotting for notebook
15 %matplotlib inline
16 plt.style.use('default')
```

### 5.2 Basic Episode

```
1 from env import GridWorldEnv
2 import numpy as np
3
4 # Create environment
5 env = GridWorldEnv(seed=42)
6
```

```

7  # Run one episode
8  state = env.reset()
9  done = False
10 total_reward = 0
11
12 while not done:
13     # Random action (replace with your policy)
14     action = np.random.randint(0, 4)
15
16     # Take action
17     next_state, reward, done, _ = env.step(action)
18     total_reward += reward
19
20     # Optional: Render
21     env.render(delay=0.1)
22
23     state = next_state
24
25 print(f"Episode finished with total reward: {total_reward}")
26 env.close()

```

### 5.3 Teacher-Student Framework

```

1  # Rendering with teacher information
2  env.render(
3      delay=0.05,
4      episode=episode,
5      learning_type="Q-learning with Teacher",
6      availability=0.7, # 70% availability
7      accuracy=0.9     # 90% accuracy
8  )

```

## 6 Common Issues and Solutions

Issue	Solution
ModuleNotFoundError: No module named 'pygame'	Install pygame: <code>pip install pygame</code>
FileNotFoundError: 'images/grid_agent.png'	Ensure the images/ folder exists in your project directory with all three PNG files
Display issues in WSL/SSH	Set environment variable: <code>export SDL_VIDEODRIVER=dummy</code> or run without rendering
Window not responding	Ensure you're calling <code>env.render()</code> in your loop; Pygame requires regular event processing
Slow training with rendering	Only render periodically (e.g., every 50-100 episodes) or increase the delay parameter

## 7 Important Notes

- Actions that would move the agent outside grid boundaries are ignored
- Actions that would move the agent into an obstacle are ignored, but the agent still receives the obstacle penalty (-10)
- The episode only terminates when the goal is reached; there is no maximum step limit in the environment itself
- Always use `seed` parameter for reproducible experiments
- The rendering is optional and can be disabled for faster training

### 7.1 Jupyter Notebook Requirements

- Submit your solution as a single `.ipynb` file
- Ensure all cells have been run and outputs are visible
- Include markdown cells to explain your approach
- Use `%matplotlib inline` for plots to display in the notebook
- The notebook should be self-contained and runnable from top to bottom
- Do not use `env.render()` in the final submission as it requires pygame display

## 8 Integration with Assignment Tasks

When implementing your reinforcement learning algorithms:

1. Initialise your Q-table with dimensions: `(11, 11, 4)`
2. Access Q-values using: `q_table[state[0], state[1], action]`
3. Track metrics as specified in the assignment requirements
4. The environment automatically handles boundaries and obstacles